

Assignment 1

Assignment 1

Biomedical Data Science

Due on Thursday 25th February 2020, 5:00pm

The assignment is marked out of 100 points, and will contribute to 20% of your final mark. Please knit this document in PDF format and submit using the gradescope link on Learn. If you can't knit to PDF directly, knit it to word and you should be able to either convert to PDF or print it and scan to PDF using a scanning app on your phone. If you have any code that doesn't run you won't be able to knit the document so comment it as you might still get some grades for partial code. Clear and reusable code will be rewarded so pay attention to indentation, choice of variable identifiers, comments, error checking, etc. An initial code chunk is provided after each subquestion but create as many chunks as you feel is necessary to make a clear report. Add plain text explanations in between the chunks as and when required and any comments necessary within code chunks to make it easier to follow your code/reasoning.

Problem 1 (25 points)

Files `longegfr1.csv` and `longegfr2.csv` (available on Learn) contain information regarding a longitudinal dataset containing records on 250 patients. For each subject, eGFR (estimated glomerular filtration rate, a measure of kidney function) was collected at irregularly spaced time points: variable "fu.years" contains the follow-up time (that is, the distance from baseline to the date when each eGFR measurement was taken, expressed in years).

Problem 1.a (4 points)

Convert the files to data tables (or tibble) and merge in an appropriate way into a single data table, then order the observations according to subject identifier and follow-up time.

```
is.nan.data.frame = function(x)
  do.call(cbind, lapply(x, is.nan))
```

Enter code here.

```
# reading datasets and setting primary keys
longegfr1 = data.table(fread('./data/longegfr1.csv'), key=c('id', 'fu.years'))
longegfr2 = data.table(fread('./data/longegfr2.csv'), key=c('ID', 'fu.years'))
colnames(longegfr2) = c('id', 'fu.years', 'egfr') # change column names to make joining easier
pk = c('id', 'fu.years')
longegfr = merge(longegfr1,
                 longegfr2,
                 by=pk,
                 all=TRUE)[order(id, fu.years)] # FULL OUTER JOIN
longegfr %>% summary()
```

```
##           id           fu.years           sex           baseline.age
##  Min.      : 1.0    Min.      :0.0000    Min.      :0.0000    Min.      :18.3
##  1st Qu.: 58.0    1st Qu.:0.8597    1st Qu.:0.0000    1st Qu.:54.7
```

```
## Median :123.0 Median :2.3682 Median :0.0000 Median :63.5
## Mean :118.9 Mean :2.6598 Mean :0.4297 Mean :63.2
## 3rd Qu.:177.0 3rd Qu.:4.4353 3rd Qu.:1.0000 3rd Qu.:74.4
## Max. :250.0 Max. :6.6283 Max. :1.0000 Max. :91.4
##
## egfr
## Min. : 4.83
## 1st Qu.: 41.05
## Median : 61.08
## Mean : 66.31
## 3rd Qu.: 86.44
## Max. :174.94
## NA's :212
```

Problem 1.b (6 points)

Compute the average eGFR and length of follow-up for each patient, then tabulate the number of patients with average eGFR in the following ranges: (0, 15], (15, 30], (30, 60], (60, 90], (90, max(eGFR)). Count and report the number of patients with missing average eGFR.

Enter code here.

```
longegfr[,mean_egfr:=mean(egfr),
        by=id]
longegfr[,count_egfr:=length(fu.years),
        by=id]
longegfr[,fu_total:=max(fu.years)-min(fu.years),
        by=id]
summary(longegfr)
```

```
##      id      fu.years      sex      baseline.age
## Min.   : 1.0   Min.   :0.0000   Min.   :0.0000   Min.   :18.3
## 1st Qu.: 58.0   1st Qu.:0.8597   1st Qu.:0.0000   1st Qu.:54.7
## Median :123.0   Median :2.3682   Median :0.0000   Median :63.5
## Mean   :118.9   Mean   :2.6598   Mean   :0.4297   Mean   :63.2
## 3rd Qu.:177.0   3rd Qu.:4.4353   3rd Qu.:1.0000   3rd Qu.:74.4
## Max.   :250.0   Max.   :6.6283   Max.   :1.0000   Max.   :91.4
##
##      egfr      mean_egfr      count_egfr      fu_total
## Min.   : 4.83   Min.   : 14.87   Min.   : 1.00   Min.   :0.000
## 1st Qu.: 41.05   1st Qu.: 44.12   1st Qu.:13.00   1st Qu.:3.817
## Median : 61.08   Median : 59.28   Median :24.00   Median :5.287
## Mean   : 66.31   Mean   : 60.53   Mean   :28.09   Mean   :4.680
## 3rd Qu.: 86.44   3rd Qu.: 76.59   3rd Qu.:39.00   3rd Qu.:6.097
## Max.   :174.94   Max.   :147.69   Max.   :88.00   Max.   :6.628
## NA's   :212     NA's   :758
```

```
longegfr
```

```
##      id fu.years sex baseline.age  egfr mean_egfr count_egfr fu_total
## 1: 1 0.0000 0 65.5 76.48 43.04333 15 6.4586
## 2: 1 0.1533 0 65.5 47.36 43.04333 15 6.4586
## 3: 1 0.6899 0 65.5 94.87 43.04333 15 6.4586
## 4: 1 1.1882 0 65.5 52.12 43.04333 15 6.4586
## 5: 1 1.8398 0 65.5 91.91 43.04333 15 6.4586
## ---
## 4027: 249 1.9521 1 50.2 91.94 75.59571 7 2.6174
```

```
## 4028: 249 2.1246 1 50.2 69.51 75.59571 7 2.6174
## 4029: 249 2.5982 1 50.2 53.28 75.59571 7 2.6174
## 4030: 249 2.6174 1 50.2 66.78 75.59571 7 2.6174
## 4031: 250 0.0000 1 48.6 101.23 101.23000 1 0.0000
```

```
longegfr$bag_egfr = cut(longegfr$mean_egfr, c(0,15,30,60,90,Inf))
```

```
longegfr_nas = longegfr[,sum(is.na(mean_egfr)), by=id]
longegfr_nas[,missing_ind:=ifelse(V1>0,1,0)]
patients_egfr_na = longegfr_nas[,sum(missing_ind)]
cat('There are', patients_egfr_na, 'patients with missing values for \'mean egfr\'')
```

```
## There are 39 patients with missing values for 'mean egfr'
```

Problem 1.c (6 points)

For patients with average eGFR in the (90,max(eGFR)) range, collect in a data table (or tibble) their identifier, sex, age at baseline, average eGFR, time of last eGFR reading and number of eGFR measurements taken.

```
# Enter code here.
```

```
egfr_90_plus = unique(longegfr[mean_egfr>90, .(id,sex,baseline.age,mean_egfr, count_egfr),])[order(mean_egfr_90_plus
```

```
##      id sex baseline.age mean_egfr count_egfr
## 1: 120  0      90.9  90.04000         2
## 2: 170  0      87.0  90.56000         2
## 3: 157  0      63.8  90.57308        13
## 4: 140  0      51.6  90.60929        28
## 5: 112  1      77.8  90.66500         6
## 6:  45  1      24.9  91.25000         1
## 7:  79  0      65.6  91.45057        35
## 8:  52  1      56.3  93.31544        57
## 9: 196  1      62.5  94.26000         2
## 10: 115  0      70.3  94.56900        10
## 11: 177  1      78.7  94.85769        26
## 12:  25  0      40.1  95.35625         8
## 13: 169  0      82.8  97.12400        10
## 14: 250  1      48.6 101.23000         1
## 15:  92  1      41.2 101.33882        17
## 16: 100  0      63.0 101.86769        13
## 17: 242  0      54.3 102.24000         4
## 18: 241  1      62.3 105.25200         5
## 19: 102  0      38.7 105.96000        10
## 20: 220  1      47.0 106.00857         7
## 21: 215  1      45.4 106.08278        54
## 22:  80  0      67.7 106.09600         5
## 23:  10  0      50.4 107.00429         7
## 24:  81  0      38.8 108.32000         8
## 25: 245  1      50.5 111.02900        10
## 26: 238  1      55.1 113.37833         6
## 27:  31  0      74.8 113.59250         8
## 28: 205  0      59.9 114.84833         6
## 29:  14  0      65.1 116.09200        10
## 30:  33  0      74.2 116.35000         4
## 31: 234  1      55.3 116.38250         4
```

```
## 32: 218 0 56.6 117.65750 4
## 33: 247 0 48.4 118.70667 9
## 34: 49 1 68.2 128.25800 5
## 35: 134 0 31.7 133.29500 2
## 36: 173 0 22.1 147.69000 1
##      id sex baseline.age mean_egfr count_egfr
```

Problem 1.d (9 points)

For patients 3, 37, 162 and 223: * Plot the patient's eGFR measurements as a function of time. * Fit a linear regression model and add the regression line to the plot. * Report the 95% confidence interval for the regression coefficients of the fitted model. * Using a different colour, plot a second regression line computed after removing the extreme eGFR values (one each of the highest and the lowest value).

The plots should be appropriately labeled and the results should be accompanied by some explanation as you would communicate it to a colleague with a medical rather than statistical background.

```
# Enter code here.
patient_id = c(3,37,162,223)
subset = longegfr[id %in% patient_id, .(id, fu.years, egfr)][order(id, fu.years)]
subset_1 = subset[id==patient_id[1]]
subset_2 = subset[id==patient_id[2]]
subset_3 = subset[id==patient_id[3]]
subset_4 = subset[id==patient_id[4]]
subset_1_rob = subset_1[!(egfr==max(egfr, na.rm=TRUE) | egfr==min(egfr, na.rm=TRUE))]
subset_2_rob = subset_2[!(egfr==max(egfr, na.rm=TRUE) | egfr==min(egfr, na.rm=TRUE))]
subset_3_rob = subset_3[!(egfr==max(egfr, na.rm=TRUE) | egfr==min(egfr, na.rm=TRUE))]
subset_4_rob = subset_4[!(egfr==max(egfr, na.rm=TRUE) | egfr==min(egfr, na.rm=TRUE))]

reg1 = lm(subset_1[,egfr] ~ subset_1[,fu.years])
reg2 = lm(subset_2[,egfr] ~ subset_2[,fu.years])
reg3 = lm(subset_3[,egfr] ~ subset_3[,fu.years])
reg4 = lm(subset_4[,egfr] ~ subset_4[,fu.years])
reg1_rob = lm(subset_1_rob[,egfr] ~ subset_1_rob[,fu.years])
reg2_rob = lm(subset_2_rob[,egfr] ~ subset_2_rob[,fu.years])
reg3_rob = lm(subset_3_rob[,egfr] ~ subset_3_rob[,fu.years])
reg4_rob = lm(subset_4_rob[,egfr] ~ subset_4_rob[,fu.years])

coef_1 = summary(reg1, conf.int=TRUE)$coefficients[2,1]
sd_1 = summary(reg1, conf.int=TRUE)$coefficients[2,2]
coef_2 = summary(reg2, conf.int=TRUE)$coefficients[2,1]
sd_2 = summary(reg2, conf.int=TRUE)$coefficients[2,2]
coef_3 = summary(reg3, conf.int=TRUE)$coefficients[2,1]
sd_3 = summary(reg3, conf.int=TRUE)$coefficients[2,2]
coef_4 = summary(reg4, conf.int=TRUE)$coefficients[2,1]
sd_4 = summary(reg4, conf.int=TRUE)$coefficients[2,2]

ci_1 = c(coef_1 - 2*sd_1 , coef_1 + 2*sd_1)
ci_2 = c(coef_2 - 2*sd_2 , coef_2 + 2*sd_2)
ci_3 = c(coef_3 - 2*sd_3 , coef_3 + 2*sd_3)
ci_4 = c(coef_4 - 2*sd_4 , coef_4 + 2*sd_4)


m = matrix(c(1,2,3,4,5,5),nrow = 3,ncol = 2,byrow = TRUE)

layout(mat = m,heights = c(0.4,0.4,0.2))
```

```

par(mar = c(2,2,1,1))
plot(subset_1[,fu.years], subset_1[,egfr], xlab='years', ylab='patient 1 egfr')
abline(reg1, col='red4')
abline(reg1_rob, col='blue4')
plot(subset_2[,fu.years], subset_2[,egfr], xlab='years', ylab='patient 2 egfr')
abline(reg2, col='red4')
abline(reg2_rob, col='blue4')
plot(subset_3[,fu.years], subset_3[,egfr], xlab='years', ylab='patient 3 egfr')
abline(reg3, col='red4')
abline(reg3_rob, col='blue4')
plot(subset_4[,fu.years], subset_4[,egfr], xlab='years', ylab='patient 4 egfr')
abline(reg4, col='red4')
abline(reg4_rob, col='blue4')
plot(1, type = "n", axes=FALSE, xlab="", ylab="")
plot_colors = c("blue", "black", "green", "orange", "pink")
legend(x = "top", inset = 0,
      legend = c("Regression Line", "Robust Regression Line"),
      col=c('red4', 'blue4'), lwd=3, cex=1.5, horiz = TRUE)

```



Problem 2 (25 points)

The MDRD4 and CKD-EPI equations are two different ways of estimating the glomerular filtration rate (eGFR) in adults: $MDRD4 = 175 \times Scr^{-1.154} \times Age^{-0.203} [\times 0.742 \text{ if female}] [\times 1.212 \text{ if black}]$, and $CKD_EPI = 141 \times \min(Scr/\kappa, 1)^\alpha \times \max(Scr/\kappa, 1)^{-1.209} \times 0.993^{Age} [\times 1.018 \text{ if female}] [\times 1.158 \text{ if black}]$, (1)


where: * Scr is serum creatinine (in mg/dL) * κ is 0.7 for females and 0.9 for males * α is -0.329 for females and -0.411 for males

Problem 2.a (7 points)

For the scr.csv dataset available on Learn, examine a summary of the distribution of serum creatinine and report the inter-quartile range. If you suspect that some serum creatinine values may have been reported in $\mu\text{mol/L}$ convert them to mg/dL by dividing by 88.42. Justify your choice of values to convert and examine the distribution of serum creatinine following any changes you have made.


Enter code here.

```
scr.dt = data.table(fread('./data/scr.csv'))  
plot(scr.dt[,scr])
```



longegfr_nas[,missing_ind:=ifelse(V1>0,1,0)]

```
scr.dt[, scr:=ifelse(scr>19, scr/88.42, scr)]  
plot(scr.dt[,scr])
```



```
IQR_scr = summary(scr.dt[,scr])[5] - summary(scr.dt[,scr])[2]
IQR_scr
```

```
## 3rd Qu.
##      1.85
```

Problem 2.b (11 points)

Compute the eGFR according to the two equations. Report (rounded to the second decimal place) mean and standard deviation of the two eGFR vectors and their Pearson correlation coefficient. Also report the same quantities according to strata of MDRD4 eGFR: 0-60, 60-90 and > 90.

Enter code here.

```
scr.dt[,scr_est1:=175*
  scr^(-1.154)*
  age^(-0.203)*
  (fifelse(scr.dt[,sex=='Female'],0.742,1, na=1))*
  (fifelse(scr.dt[,ethnic=='Black'],1.212,1,na=1))]

scr.dt[,min_scr_1:=pmin(scr/fifelse(scr.dt[,sex=='Female'],0.7,0.9,na=1),1, na.rm=TRUE)^(fifelse(scr.dt
scr.dt[,max_scr_1:=pmax(scr/fifelse(scr.dt[,sex=='Female'],0.7,0.9,na=1),1, na.rm=TRUE)^(-1.209)]
scr.dt[,scr_est2:=141*
  min_scr_1*
  max_scr_1*
  0.993^age*
  (fifelse(scr.dt[,sex=='Female'],1.018,1, na=1))*
  (fifelse(scr.dt[,ethnic=='Black'],1.158,1,na=1))]

scr.dt
```

```
##      age scr    sex ethnic  scr_est1 min_scr_1 max_scr_1  scr_est2
##  1:  48 1.2 Female  Other   47.94848  1.000000  0.5211868  53.39791
##  2:   7 0.8  Male   Black  184.85020  1.049600  1.0000000  163.15339
##  3:  62 1.8 Female  <NA>   28.51031  1.000000  0.3192266   29.64281
```

```
## 4: 48 3.8 Female Other 12.67885 1.000000 0.1293496 13.25244
## 5: 51 1.4 Male Other 53.42808 1.000000 0.5861522 57.76186
## ---
## 397: 55 0.5 Female Other 128.09641 1.117059 1.0000000 108.95613
## 398: 42 1.2 Male Black 80.47221 1.000000 0.7062347 85.85096
## 399: 12 0.6 Male Other 190.53629 1.181336 1.0000000 153.10297
## 400: 17 1.0 Female Black 88.54488 1.000000 0.6497160 95.83766
## 401: 58 1.1 Female Other 51.01512 1.000000 0.5790016 55.29719

mean_est_1 = round(mean(scr.dt$scr_est1, na.rm=TRUE),2)
sd_est_1 = round(sd (scr.dt$scr_est1, na.rm=TRUE),2)
mean_est_2 = round(mean(scr.dt$scr_est2, na.rm=TRUE),2)
sd_est_2 = round(sd (scr.dt$scr_est2, na.rm=TRUE),2)

mean_est_1

## [1] 61.11

sd_est_1

## [1] 49.82

mean_est_2

## [1] 61.77

sd_est_2

## [1] 42.44

cor(scr.dt$scr_est1,scr.dt$scr_est2, method='pearson', use = "complete.obs")


## [1] 0.9531475
```

Problem 2.c (7 points)

Produce a scatter plot of the two eGFR vectors, and add vertical and horizontal lines (i.e.) corresponding to median, first and third quartiles. Is the relationship between the two eGFR equations linear? Justify your answer.

```
# Enter code here.
plot(scr.dt$scr_est1,scr.dt$scr_est2, xlab='MDRD4', ylab='CKD-EPI', main='Linearity check between MDRD4
abline(v=mean_est_1, col='blue4')
abline(v=quantile(scr.dt$scr_est1, na.rm=TRUE)[2], col='blue4')
abline(v=quantile(scr.dt$scr_est1, na.rm=TRUE)[4], col='blue4')
abline(h=mean_est_2, col='red4')
abline(h=quantile(scr.dt$scr_est2, na.rm=TRUE)[2], col='red4')
abline(h=quantile(scr.dt$scr_est2, na.rm=TRUE)[4], col='red4')
```


Linearity check between MDRD4 & CKD-EPI



Problem 3 (31 points)

You have been provided with electronic health record data from a study cohort. Three CSV (Comma Separated Variable) files are provided on learn.

The first file is a cohort description file `cohort.csv` file with fields: * `id` = study identifier * `yob` = year of birth * `age` = age at measurement * `bp` = systolic blood pressure * `albumin` = last known albuminuric status (categorical) * `diabetes` = diabetes status

The second file `lab1.csv` is provided by a laboratory after measuring various biochemistry levels in the cohort blood samples. Notice that a separate lab identifier is used to anonymise results from the cohort. The year of birth is also provided as a check that the year of birth aligns between the two merged sets. * `LABID` = lab identifier * `yob` = year of birth * `urea` = blood urea * `creatinine` = serum creatinine * `glucose` = random blood glucose

To link the two data files together, a third linker file `linker.csv` is provided. The linker file includes a `LABID` identifier and the corresponding cohort `id` for each person in the cohort.

Problem 3.a (6 points)

Using all three files provided on learn, load and merge to create a single data table based dataset `cohort.dt`. This will be used in your analysis. Perform assertion checks to ensure that all identifiers in `cohort.csv` have been accounted for in the final table and that any validation fields are consistent between sets. After the checks are complete, drop the identifier that originated from lab dataset `LABID`. Ensure that a single `yob` field remains and rename it. Ensure that the `albumin` field is converted to a factor and the ordering of the factor is 1="normo", 2="micro", 3="macro".

```
# Enter code here.
cohort = data.table(fread('./data/cohort.csv', stringsAsFactors = TRUE))
lab1 = data.table(fread('./data/lab1.csv'))
linker = data.table(fread('./data/linker.csv'))
tmp = merge(cohort, linker, by=c('id'), all=TRUE)
```

```
cohort.dt = merge(tmp, lab1, by.x=c('LABID', 'yob'), by.y=c('LABID', 'yob'), all=TRUE)
cohort.dt$albumin = ordered(cohort.dt$albumin,
                             levels = c('normo','micro','macro'))
cohort.dt
```

```
##      LABID  yob      id age  bp diabetes albumin urea creatinine glucose
##  1:  LID_1 1986 PID_285 33  80         0  normo 37.0   106.104    100
##  2:  LID_10 1980 PID_153 39  70         1  normo 20.0    70.736    121
##  3: LID_100 1951  PID_13  68  70         1  micro 72.0   185.682    208
##  4: LID_101 1965 PID_110 54  70         1   <NA> 50.1   167.998    233
##  5: LID_102 1953 PID_222 66  70         1  micro 30.0   150.314    248
## ---
## 396: LID_95 1962 PID_254 57  80         0  normo 17.0   106.104    119
## 397: LID_96 1978 PID_297 41  70         0  normo 38.0    53.052    125
## 398: LID_97 1964 PID_119 55  70         0  micro 25.0   106.104     99
## 399: LID_98 1974 PID_236 45  70         0  micro 93.0   203.366    113
## 400: LID_99 1963 PID_100 56 180         1  normo 24.0   106.104    298
```

Problem 3.b (10 points)

Create a copy of the dataset where you will impute all missing values. Update any missing age fields using the year of birth, for all other continuous variables write a function called `impute.to.mean` and impute to mean, impute any categorical variable to the mode. Compare the distributions of the imputed and non-imputed variables and decide which ones to keep for further analysis. Justify your answer.

Enter code here.

```
cohort_complete = copy(cohort.dt)
cohort_complete[is.na(age), age:=abs(year(Sys.Date())-round(yob))]
cohort_complete
```

```
##      LABID  yob      id age  bp diabetes albumin urea creatinine glucose
##  1:  LID_1 1986 PID_285 33  80         0  normo 37.0   106.104    100
##  2:  LID_10 1980 PID_153 39  70         1  normo 20.0    70.736    121
##  3: LID_100 1951  PID_13  68  70         1  micro 72.0   185.682    208
##  4: LID_101 1965 PID_110 54  70         1   <NA> 50.1   167.998    233
##  5: LID_102 1953 PID_222 66  70         1  micro 30.0   150.314    248
## ---
## 396: LID_95 1962 PID_254 57  80         0  normo 17.0   106.104    119
## 397: LID_96 1978 PID_297 41  70         0  normo 38.0    53.052    125
## 398: LID_97 1964 PID_119 55  70         0  micro 25.0   106.104     99
## 399: LID_98 1974 PID_236 45  70         0  micro 93.0   203.366    113
## 400: LID_99 1963 PID_100 56 180         1  normo 24.0   106.104    298
```

```
#' This is a function that takes as impute a column of a data.table and imputes
#' the NAs with its mean / mode if the vector is numeric or categorical respectively.
#' @param x A vector of numeric or categorical values for which the NAs will be imputed.
impute.to.mean = function(x) {
  if (is.numeric(x)){
    if (all(na.omit(x) %in% 0:1)){
      x[is.na(x)] = unique(x)[which.max(tabulate(match(x, unique(x))))]
    } else {x[is.na(x)] = mean(x, na.rm=TRUE)}
  } else if (is.factor(x)){x[is.na(x)] = unique(x)[which.max(tabulate(match(x, unique(x))))]}
  return(x)
}
```

```
# diab.dt.imputed2 = diab.dt %>% copy() %>%
#           .[, (numcols) := lapply(.SD, impute.to.median), .SDcols = numcols]
# bp,diabetes,urea,creatinine,glucose, albumin
numcols = cohort_complete %>% select(bp,diabetes,urea,creatinine,glucose, albumin) %>% colnames
cohort_complete %>% .[, (numcols) := lapply(.SD, impute.to.mean), .SDcols = numcols]
unique(cohort.dt$bp)[which.max(tabulate(match(cohort.dt$bp, unique(cohort.dt$bp))))]
```

```
## [1] 80
```

```
summary(cohort_complete)
```

```
##      LABID          yob          id          age
## Length:400      Min.   :1929 Length:400      Min.   : 2.00
## Class :character 1st Qu.:1955 Class :character 1st Qu.:42.00
## Mode  :character Median :1965 Mode  :character Median :54.00
##              Mean  :1968              Mean  :51.52
##              3rd Qu.:1977              3rd Qu.:64.00
##              Max.   :2017              Max.   :90.00
##      bp      diabetes      albumin      urea
## Min.   : 50.00      Min.   :0.0000      normo:245      Min.   : 1.50
## 1st Qu.: 70.00      1st Qu.:0.0000      micro:130      1st Qu.: 27.00
## Median : 78.23      Median :0.0000      macro: 25      Median : 44.00
## Mean   : 76.47      Mean   :0.3375              Mean   : 57.43
## 3rd Qu.: 80.00      3rd Qu.:1.0000              3rd Qu.: 61.75
## Max.   :180.00      Max.   :1.0000              Max.   :391.00
##      creatinine      glucose
## Min.   : 35.37      Min.   : 22
## 1st Qu.: 79.58      1st Qu.:101
## Median :123.79      Median :126
## Mean   :271.67      Mean   :148
## 3rd Qu.:271.67      3rd Qu.:150
## Max.   :6719.92      Max.   :490
```

```
cohort_complete
```

```
##      LABID yob      id age bp diabetes albumin urea creatinine glucose
## 1:  LID_1 1986 PID_285 33 80      0  normo 37.0  106.104  100
## 2:  LID_10 1980 PID_153 39 70      1  normo 20.0   70.736  121
## 3:  LID_100 1951 PID_13 68 70      1  micro 72.0  185.682  208
## 4:  LID_101 1965 PID_110 54 70      1  normo 50.1  167.998  233
## 5:  LID_102 1953 PID_222 66 70      1  micro 30.0  150.314  248
## ---
## 396:  LID_95 1962 PID_254 57 80      0  normo 17.0  106.104  119
## 397:  LID_96 1978 PID_297 41 70      0  normo 38.0   53.052  125
## 398:  LID_97 1964 PID_119 55 70      0  micro 25.0  106.104   99
## 399:  LID_98 1974 PID_236 45 70      0  micro 93.0  203.366  113
## 400:  LID_99 1963 PID_100 56 180      1  normo 24.0  106.104  298
```

```
m = matrix(c(1,2,3,4,5,6,7,7),nrow = 3,ncol = 3,byrow = TRUE)
```

```
layout(mat = m,heights = c(0.4,0.4,0.2))
```

```
par(mar = c(2,2,1,1))
```

```
barplot(rbind(table(cohort_complete$bp), table(cohort.dt$bp)), beside=TRUE, col=c('red4','blue4'), main=
```

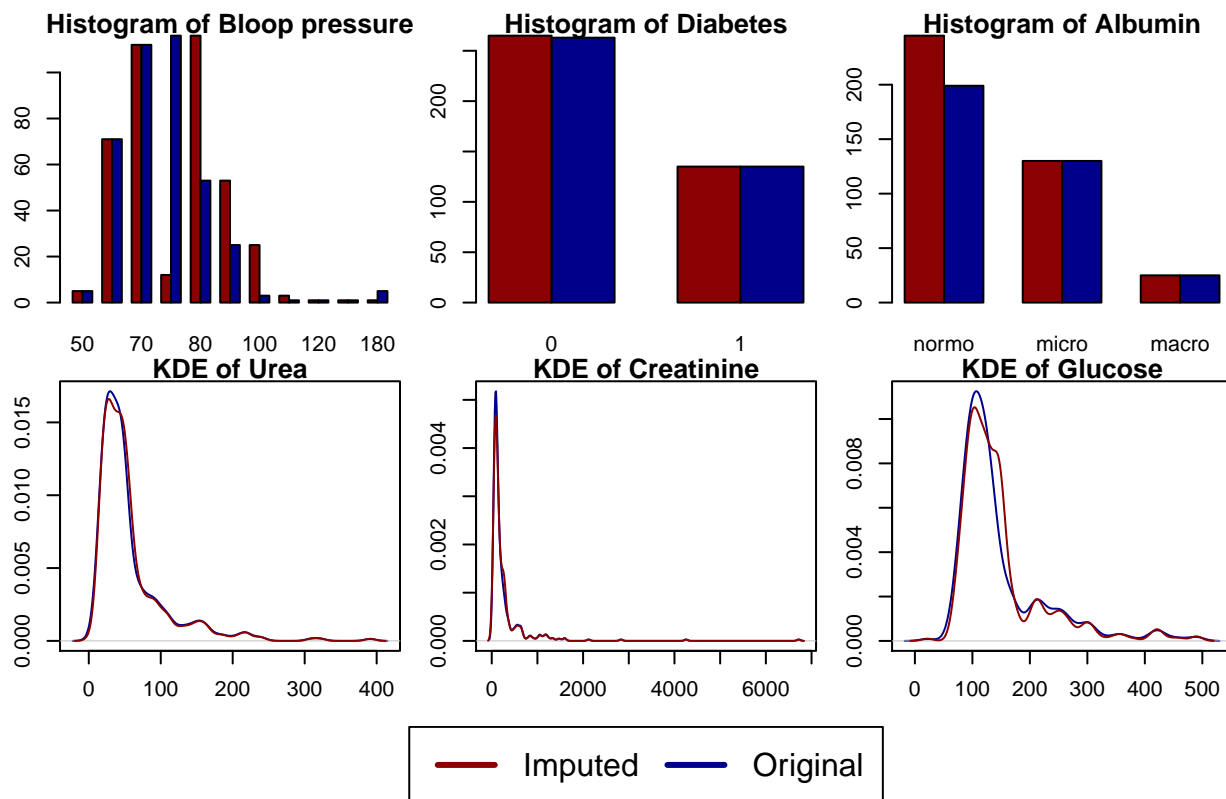
```
## Warning in rbind(table(cohort_complete$bp), table(cohort.dt$bp)): number of
```

```
## columns of result is not a multiple of vector length (arg 2)
```

```

barplot(rbind(table(cohort_complete$diabetes), table(cohort.dt$diabetes)), beside=TRUE, col=c('red4','blue4'))
barplot(rbind(table(cohort_complete$albumin), table(cohort.dt$albumin)), beside=TRUE, col=c('red4','blue4'))
plot(density(cohort.dt$urea, na.rm=TRUE), col='blue4', main='KDE of Urea')
lines(density(cohort_complete$urea), col='red4')
plot(density(cohort.dt$creatinine, na.rm=TRUE), col='blue4', main='KDE of Creatinine')
lines(density(cohort_complete$creatinine), col='red4')
plot(density(cohort.dt$glucose, na.rm=TRUE), col='blue4', main='KDE of Glucose')
lines(density(cohort_complete$glucose), col='red4')
plot(1, type = "n", axes=FALSE, xlab="", ylab="")
legend(x = "top", inset = 0,
       legend = c("Imputed", "Original"),
       col=c('red4', 'blue4'), lwd=3, cex=1.5, horiz = TRUE)

```



Problem 3.c (6 points)

Plot boxplots of potential predictors for diabetes grouped by cases and controls and use these to decide which predictors to keep for future analysis. For any categorical variables create a table instead. Justify your answers.

```

# Enter code here.
# bp,diabetes,urea,creatinine,glucose, albumin
bp = cohort_complete$bp
urea = cohort_complete$urea
creatinine = cohort_complete$creatinine
glucose = cohort_complete$glucose
albumin = cohort_complete$albumin
diabetes = cohort_complete$diabetes


m = matrix(c(1,2,3,4,5,3),nrow = 2,ncol = 3,byrow = TRUE)

```

```

layout(mat = m,heights = c(0.4,0.4,0.2))
par(mar = c(2,2,1,1))
boxplot(bp ~ diabetes, data=cohort_complete, main="Blood press stratified by Diabetes")
boxplot(urea ~ diabetes, data=cohort_complete, main="Urea stratified by Diabetes")
barplot(rbind(table(cohort_complete$diabetes), table(cohort.dt$albumin)), beside=TRUE, col=c('red4','blue4'))
boxplot(creatinine ~ diabetes, data=cohort_complete, main="Creatine stratified by Diabetes")
boxplot(glucose ~ diabetes, data=cohort_complete, main="Glucose stratified by Diabetes")

```



```

albumin_table = table(albumin, diabetes)
colnames(albumin_table) = c("No Diabetes", "Diabetes")
albumin_table

```

```

##      diabetes
## albumin No Diabetes Diabetes
##  normo      192      53
##  micro      61      69
##  macro      12      13

```

Problem 3.d (9 points)

Use your findings from the previous exercise fit an appropriate model of diabetes with two predictors. Print a summary and explain the results as you would communicate it to a colleague with a medical rather than statistical background.

```

# Enter code here.
model1 = glm(diabetes ~ glucose + albumin, family='binomial')
summary(model1)

```

```
##
## Call:
## glm(formula = diabetes ~ glucose + albumin, family = "binomial")
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.0990  -0.7013  -0.5208   0.6802   2.2565
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -3.151621    0.416888  -7.560 4.03e-14 ***
## glucose      0.017834    0.002513   7.095 1.29e-12 ***
## albumin.L     0.534820    0.342585   1.561  0.1185
## albumin.Q    -0.448139    0.256400  -1.748  0.0805 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 511.49  on 399  degrees of freedom
## Residual deviance: 389.59  on 396  degrees of freedom
## AIC: 397.59
##
## Number of Fisher Scoring iterations: 5
```

Problem 4 (19 points)

Problem 4.a. (9 points)

Add a third predictor to the final model from problem 3, perform a likelihood ratio test to compare both models and report the p-value for the test. Is there any support for the additional term? Plot a ROC curve for both models and report the AUC, explain the results as you would communicate it to a colleague with a medical rather than statistical background.

```
# Enter code here.
model2 = glm(diabetes ~ glucose + albumin + cohort_complete$age, family='binomial')
summary(model2)
```

```
##
## Call:
## glm(formula = diabetes ~ glucose + albumin + cohort_complete$age,
##      family = "binomial")
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.7615  -0.7072  -0.4022   0.6420   2.5238
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -5.451612    0.681802  -7.996 1.29e-15 ***
## glucose         0.016281    0.002538   6.416 1.40e-10 ***
## albumin.L       0.642553    0.373277   1.721  0.0852 .
## albumin.Q      -0.247191    0.277030  -0.892  0.3722
## cohort_complete$age 0.047137    0.009543   4.940 7.83e-07 ***
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## (Dispersion parameter for binomial family taken to be 1)
```

```
##
```

```
##      Null deviance: 511.49  on 399  degrees of freedom
```

```
## Residual deviance: 359.43  on 395  degrees of freedom
```

```
## AIC: 369.43
```

```
##
```

```
## Number of Fisher Scoring iterations: 5
```

```
pval = pchisq(model1$deviance - model2$deviance, df=1, lower.tail=FALSE)
```

```
signif(pval, 2)
```

```
## [1] 4e-08
```

```
roc1 = roc(cohort_complete$diabetes, model1$fitted.values , plot=TRUE, xlim = c(0,1))
```


```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
roc2 = roc(cohort_complete$diabetes, model2$fitted.values , plot=TRUE, xlim = c(0,1), add=TRUE, col="red")
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```



```
roc1$auc
```

```
## Area under the curve: 0.8089
```

```
roc2$auc
```

```
## Area under the curve: 0.8522
```

Problem 4.b (10 points)

Perform 10-folds cross-validation for your chosen model and report the mean cross-validated AUCs.

```
# Enter code here.
set.seed(1903)
k = 10
folds = createFolds(cohort_complete$diabetes, k=k)

pred.cv <- NULL
regr.cv <- NULL
auc.cv <- numeric(k)
for(f in 1:k) {
  train.idx <- setdiff(1:nrow(cohort_complete), folds[[f]])
  regr.cv[[f]] <- glm(diabetes ~ glucose + albumin + cohort_complete$age, family='binomial')
  test.idx <- folds[[f]]
  pred.cv[[f]] <- data.frame(obs = cohort_complete$diabetes[test.idx],
                             pred = predict(regr.cv[[f]], newdata = cohort_complete, type = "response"))
  auc.cv[f] <- roc(obs ~ pred, data = pred.cv[[f]])$auc
}

round(mean(auc.cv), 3)

## [1] 0.843
```